




Contoso Engineering (illustrative mock)

AI Maturity Assessment summary

AI-Assisted SDLC maturity: 9 dimensions, 61 questions

ASSESSMENT DATE	2026-09-27
FRAMEWORK VERSION	2.0.1
RESPONDENTS	14
GENERATED ON	2026-09-29

1. Overall result



Coverage is OK with 37 or more questions answered, WARNING from 25, BLOCKED below 25.

1.1 Flags

Amplification risk: D8 (L1). These foundations are at least one band below the overall level (L2). DORA finds that AI amplifies existing strengths and weaknesses [3].
Perception gap: D1 (3.45 / 2.50), D2 (2.69 / 2.02), D3 (2.58 / 1.80), D5 (2.60 / 1.80), D6 (2.38 / 1.59), D7 (2.33 / 1.45), D8 (1.91 / 0.79), D9 (2.10 / 1.31). Values are executive / hands-on scores. Self-reported impact can differ from measured impact [35]; confirm with telemetry.
Low confidence: D6. More than 30% of the answers are NA.
Unverified L3/L4: D1-Q1, D1-Q3, D1-Q4, D1-Q5, D1-Q7, D4-Q1, D4-Q3, D4-Q4, D4-Q5, D4-Q6, D4-Q7. Scored L3 or higher, but fewer than half of the answers carry evidence. Treat as unverified until confirmed.

Scores by dimension

ID	DIMENSION	SCORE	LEVEL	ANSWERED	NA	PRIORITY
D1	AI Strategy, Policy and Governance	2.82	L3 Scaling	7 / 7	5%	P3 Low
D2	Enablement, Skills and Culture	2.28	L2 Adopting	6 / 6	2%	P3 Low
D3	Plan, Specify and Design	2.09	L2 Adopting	6 / 6	6%	P2 Medium
D4	Code and Context Engineering	2.72	L3 Scaling	8 / 8	7%	P3 Low
D5	Review, Quality and Testing	2.10	L2 Adopting	7 / 7	5%	P2 Medium
D6	Security and AI Supply Chain low confidence	1.90	L2 Adopting	7 / 7	42%	P2 Medium
D7	Deliver and Operate	1.75	L2 Adopting	6 / 6	5%	P2 Medium
D8	Engineering Foundations (AI amplifiers)	1.09	L1 Exploring	7 / 7	4%	P1 High
D9	Measurement, Value and AI FinOps	1.61	L2 Adopting	7 / 7	6%	P2 Medium

2.1 Dimensions by persona

Dimension	All	Engineering manager / tech lead n=2 · low sample	Executive (CTO, VP, Director) n=3	Platform / DevOps / SRE engineer n=2 · low sample	Product / program manager n=1 · low sample	Security / AppSec n=1 · low sample	Software engineer / developer n=5
D1	2.82	2.86	3.45	2.79	3.29	2.43	2.38
D2	2.28	2.33	2.69	2.50	3.00	2.00	1.83
D3	2.09	2.17	2.58	1.83	2.80	1.67	1.82
D4	2.72	2.94	3.17	2.44	3.00	2.40	2.51
D5	2.10	2.43	2.60	1.93	2.33	1.43	1.86
D6	1.90	2.29	2.38	1.71	2.00	1.25	1.58
D7	1.75	1.92	2.33	1.50	2.33	1.50	1.42
D8	1.09	1.14	1.91	0.86	1.43	1.00	0.72
D9	1.61	1.79	2.10	1.36	2.00	1.00	1.36

Persona is the role in R-Q1. Groups with fewer than 3 respondents are marked as low sample.

2.2 Evidence cross-checks

Repository scan (D4-Q4, D4-Q5)

12 repositories: 3 unconfigured, 5 grounded, 3 agent and 1 orchestration on the RAMP levels [47] (pattern-based scan). 75% have committed AI instructions or configuration, so repository coverage supports at most L3 for D4-Q4. 4 of 8 instruction files were committed once and never changed. 33% of the repositories have custom agents, prompt files or skills (context for D4-Q5).

Assessment: D4-Q4 = 2.50 (L3). Assessment: D4-Q5 = 2.79 (L3). Consistent with the answers.

Copilot usage metrics (D4-Q1, D9-Q1)

Report day 2026-09-27: 140 users in the 28-day window; 71% Code first or higher, 21% Agent first or higher and 4% Multi-agent [6]. By the agentic signal of the scale, this supports at most L2 for D4-Q1. 60% of the 200 licensed seats were active in the last 28 days. A usage metrics export exists, which is evidence for D9-Q1 (adoption tracked beyond seat counts).

Assessment: D4-Q1 = 2.62 (L3). Assessment: D9-Q1 = 1.79 (L2). **The answers for D4-Q1 are above what the evidence supports; review them with the evidence before presenting.**

2.3 Developer Survey context

From developer-survey-maturity-2026-09-29.json (5 respondents). Context only: survey results never change v2 scores, and the survey uses the v1 bands, so compare scores rather than level names.

SURVEY DIMENSION	TEAM SCORE	RELATED V2 QUESTIONS (ASSESSMENT SCORE)
DS-D2 Copilot Adoption	0.80 L1 Developing	D4-Q1 (2.62), D4-Q2 (2.64), D9-Q1 (1.79)
DS-D3 MS/GH Tooling Breadth	2.40 L2 Defined	D4-Q3 (2.67), D4-Q6 (2.79), D3-Q2 (2.00), D6-Q1 (2.00)
DS-D4 AI Dev Practices	2.68 L3 Managed	D3-Q2 (2.00), D5-Q5 (2.31), D2-Q6 (2.15)
DS-D5 Agent Concepts Mastery	2.56 L3 Managed	D2-Q4 (2.21), D4-Q5 (2.79)
DS-D6 Instructions Maturity	2.31 L2 Defined	D4-Q4 (2.50), D4-Q5 (2.79)
DS-D7 Best Practices	2.91 L3 Managed	D2-Q2 (2.36), D9-Q4 (1.57)
DS-D8 Security & Governance	1.92 L2 Defined	D1-Q2 (2.93), D6-Q1 (2.00), D6-Q4 (1.71), D6-Q5 (1.50), D6-Q7 (2.08)

3. Starting backlog

The five lowest-scoring questions, with the L3 anchor as the first target. Confirm each one with evidence before planning work.

ID	QUESTION	SCORE	L3 TARGET	BASIS
D8-Q7	Automated testing as a control system	0.75	Most services have layered automated tests that run on every PR within agreed time budgets.	[3]
D8-Q4	Healthy data ecosystem	0.92	Key data is cataloged with owners and quality indicators; most questions can be answered within an hour.	[1][2]
D8-Q5	Reproducible environments for humans and agents	1.00	Most repositories define a reproducible environment; agents use the same definition.	[8]
D8-Q1	Version control for everything	1.07	All five asset types are versioned for most services.	[1][2]
D8-Q2	Commit frequency and fast rollback	1.23	Most engineers commit at least daily; reverting a change is routine and fast.	[2]

3.1 Recommended strategies

#	STRATEGY	CUMULATIVE PRIORITY	PRIORITY	HORIZON	FIRST ACTION	BASIS
1	S2 Foundry + SRE D8, D9, D7	4.55	P1 High	Next quarter	Start with D8-Q7 (Automated testing as a control system). Target the L3 anchor: Most services have layered automated tests that run on every PR within agreed time budgets.	[1][2] [3][8]
2	S5 GitHub Copilot Acceleration D9, D3, D5, D2, D4, D1	4.39	P2 Medium	Semester	Start with D9-Q3 (Pull request flow metrics). Target the L3 anchor: PR lifecycle metrics are reported per organization; agent-authored PRs are identified separately, including their post-merge fix rate.	[2][6] [18][28] [32][34] [44][50]
3	S3 App Modernization D8, D3	2.83	P1 High	Next quarter	Start with D8-Q7 (Automated testing as a control system). Target the L3 anchor: Most services have layered automated tests that run on every PR within agreed time budgets.	[1][2] [3][8]
4	S6 Agentic Activation D7, D6, D4, D1	2.81	P2 Medium	Semester	Start with D6-Q5 (AI supply chain). Target the L3 anchor: A review process covers AI components; SBOMs and build provenance are produced for most builds.	[19][32] [38][39] [40][43] [45][52]
5	S7 Security & Governance D6, D5, D1	2.19	P2 Medium	Semester	Start with D6-Q5 (AI supply chain). Target the L3 anchor: A review process covers AI components; SBOMs and build provenance are produced for most builds.	[7][19] [38][39] [40][43] [52]
6	S1 GitHub Migration D8	1.91	P1 High	Next quarter	Start with D8-Q7 (Automated testing as a control system). Target the L3 anchor: Most services have layered automated tests that run on every PR within agreed time budgets.	[1][2] [3][8]

Not recommended now: S4 (monitor (only low-priority gaps))

4. Method

LEVEL	SCORE BAND
L0 Not started	[0.0, 0.8)
L1 Exploring	[0.8, 1.6)
L2 Adopting	[1.6, 2.4)
L3 Scaling	[2.4, 3.2)
L4 AI-native	[3.2, 4.0]

Question score: mean of respondent values (NA excluded). Dimension score: mean of its question scores. Overall: weighted mean of dimensions (equal weights by default). Gap = target (3.0 by default) minus score; priority = weight × gap: P0 from 2.4, P1 from 1.6, P2 from 0.9.

Answers are perceptions. Triangulate D4, D5 and D9 with telemetry and evidence before presenting results. Thresholds and weights are kit design choices, not an industry standard.

5. References

- [1] DORA. DORA AI Capabilities Model. Google Cloud, 2025. <https://dora.dev/ai/capabilities-model/report/>
- [2] DORA. DORA AI Capabilities Model: Survey Questions. Google Cloud, 2025. <https://dora.dev/ai/capabilities-model/questions/>
- [3] Google Cloud. Announcing the 2025 DORA Report: State of AI-Assisted Software Development. 2025. <https://cloud.google.com/blog/products/ai-machine-learning/announcing-the-2025-dora-report>
- [6] GitHub Docs. GitHub Copilot usage metrics. <https://docs.github.com/en/copilot/concepts/billing-and-usage/copilot-usage-metrics/copilot-metrics>
- [7] GitHub Docs. Risks and mitigations for GitHub Copilot cloud agent. <https://docs.github.com/en/enterprise-cloud@latest/copilot/concepts/security-governance-and-network-settings/risks-and-mitigations>
- [8] GitHub Docs. Best practices for using GitHub Copilot to work on tasks. <https://docs.github.com/en/enterprise-cloud@latest/copilot/tutorials/cloud-agent/get-the-best-results>
- [12] GitHub Docs. About GitHub Copilot code review. <https://docs.github.com/en/copilot/concepts/agents/code-review>
- [14] GitHub Docs. Application card: GitHub Copilot Agents. <https://docs.github.com/en/copilot/responsible-use/agents>
- [15] GitHub Docs. Managing policies and features for GitHub Copilot in your organization. <https://docs.github.com/en/copilot/how-tos/administer-copilot/manage-for-organization/manage-policies>
- [16] Microsoft Azure Blog. Agentic DevOps: Evolving software development with GitHub Copilot and Microsoft Azure. 2025. <https://azure.microsoft.com/en-us/blog/agentic-devops-evolving-software-development-with-github-copilot-and-microsoft-azure/>
- [17] Microsoft for Developers. Agentic DevOps in action: Reimagining every phase of the developer lifecycle. 2025. <https://developer.microsoft.com/blog/reimagining-every-phase-of-the-developer-lifecycle/>
- [18] Microsoft Learn. AI strategy: Cloud Adoption Framework. <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/ai/strategy>
- [19] Microsoft Learn. Govern and secure AI agents: Cloud Adoption Framework. <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/ai-agents/governance-security-across-organization>
- [20] Microsoft Learn. Responsible AI policies: Cloud Adoption Framework. <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/ai/responsible-ai-policies>
- [21] Microsoft. Responsible AI Principles and Approach. <https://www.microsoft.com/en-us/ai/principles-and-approach>
- [22] Microsoft Azure Blog. Announcing Azure Copilot agents and AI infrastructure innovations. 2025. <https://azure.microsoft.com/en-us/blog/announcing-azure-copilot-agents-and-ai-infrastructure-innovations/>
- [23] Anthropic. Anthropic Economic Index: AI's impact on software development. 2025-04-28. <https://www.anthropic.com/research/impact-software-development>
- [24] Anthropic. How Claude Code is used in practice. 2026-06-16. <https://www.anthropic.com/research/claude-code-expertise>
- [25] Anthropic. Effective context engineering for AI agents. 2025-09-29. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>
- [26] Anthropic. Building effective agents. 2024-12-19. <https://www.anthropic.com/engineering/building-effective-agents>
- [27] Anthropic. Best practices for Claude Code. Claude Code Docs. <https://code.claude.com/docs/en/best-practices>
- [28] Gartner. Gartner Says 75% of Enterprise Software Engineers Will Use AI Code Assistants by 2028. 2024-04-11. <https://www.gartner.com/en/newsroom/press-releases/2024-04-11-gartner-says-75-percent-of-enterprise-software-engineers-will-use-ai-code-assistants-by-2028>
- [30] Gartner. Gartner Says Generative AI will Require 80% of Engineering Workforce to Upskill Through 2027. 2024-10-03. <https://www.gartner.com/en/newsroom/press-releases/2024-10-03-gartner-says-generative-ai-will-require-80-percent-of-engineering-workforce-to-upskill-through-2027>

- [31] Gartner. Gartner Says the Market for Enterprise AI Coding Agents Is Entering a New Phase of Expansion and Competitive Realignment. 2026-05-20. <https://www.gartner.com/en/newsroom/press-releases/2026-05-20-gartner-says-the-market-for-enterprise-ai-coding-agents-is-entering-a-new-phase-of-expansion-and-competitive-realignment>
- [32] Gartner. Gartner Predicts AI Coding Costs Will Surpass Average Developer's Salary by 2028 as Token Consumption Surges. 2026-06-24. <https://www.gartner.com/en/newsroom/press-releases/2026-06-24-gartner-predicts-ai-coding-costs-will-surpass-average-developer-salary-by-2028-as-token-consumption-surges>
- [34] Cui, Z. K., Demirer, M., Jaffe, S., Musolff, L., Peng, S., Salz, T. The Effects of Generative AI on High-Skilled Work: Evidence from Three Field Experiments with Software Developers. *Management Science*, published online 2026-02-27. <https://doi.org/10.1287/mnsc.2025.00535> (pre-print: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4945566)
- [35] METR. Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity. 2025-07-10. <https://metr.org/blog/2025-07-10-early-2025-ai-experienced-os-dev-study/> (paper: <https://arxiv.org/abs/2507.09089>)
- [37] Stack Overflow. 2025 Developer Survey: AI. <https://survey.stackoverflow.co/2025/ai>
- [38] OWASP GenAI Security Project. OWASP Top 10 for LLM Applications 2025. <https://genai.owasp.org/resource/owasp-top-10-for-llm-applications-2025/>
- [39] OWASP GenAI Security Project. OWASP Top 10 for Agentic Applications for 2026. 2025-12-09. <https://genai.owasp.org/resource/owasp-top-10-for-agentic-applications-for-2026/>
- [40] NIST. SP 800-218A: Secure Software Development Practices for Generative AI and Dual-Use Foundation Models: An SSDF Community Profile. 2024-07. <https://csrc.nist.gov/pubs/sp/800/218/a/final>
- [41] NIST. AI Risk Management Framework. <https://www.nist.gov/itl/ai-risk-management-framework>
- [42] ISO. ISO/IEC 42001:2023: AI management systems. <https://www.iso.org/standard/42001>
- [43] OpenSSF. SLSA: Supply-chain Levels for Software Artifacts. <https://slsa.dev/>
- [44] Forsgren, N., Storey, M.-A., Maddila, C., Zimmermann, T., Houck, B., Butler, J. The SPACE of Developer Productivity. *ACM Queue* 19(1), 2021-03-06. <https://queue.acm.org/detail.cfm?id=3454124> (DOI: <https://doi.org/10.1145/3454122.3454124>)
- [45] Liu, B., Qiu, H., Goiri, Í., Fonseca, R., Bianchini, R., Choukse, E. Agentic Coding in the Wild: Characterizing GitHub Copilot Traces at Production Scale. *arXiv:2608.00101*, 2026-07-30. <https://arxiv.org/abs/2608.00101>
- [46] Arabat, A., Sayagh, M. Toward Instructions-as-Code: Understanding the Impact of Instruction Files on Agentic Pull Requests. *arXiv:2606.13449*, 2026-06-11. <https://arxiv.org/abs/2606.13449>
- [47] Denisov-Blanch, Y., Agarwal, S., Azaletskiy, P., He, H., Schaeffer, R., Miranda, B., Vasilescu, B., Koyejo, S. A Few Pages of Markdown: Committed AI Configuration and Lower Quality Cost after Coding-Agent Adoption. *arXiv:2608.25241*, 2026-08-26. <https://arxiv.org/abs/2608.25241>
- [48] Shen, J. H., Tamkin, A. How AI Impacts Skill Formation. *arXiv:2601.20245*, 2026-01-28. <https://arxiv.org/abs/2601.20245> (Anthropic summary: <https://www.anthropic.com/research/AI-assistance-coding-skills>)
- [49] Pinna, G., Gong, J., Williams, D., Sarro, F. Comparing AI Coding Agents: A Task-Stratified Analysis of Pull Request Acceptance. *arXiv:2602.08915*, 2026-02-09. <https://arxiv.org/abs/2602.08915>
- [50] Takerngsaksiri, W., Duong, N., Barnett, S. Who Finishes the Job? A Study of Follow-Up Fixes and Commit Authorship on AI Coding Agent Pull Requests. *arXiv:2609.26847*, 2026-09-22. <https://arxiv.org/abs/2609.26847>
- [52] Sakib, A. H. M. N., Banik, D., Jadliwala, M. Trust but Verify? Uncovering the Security Debt of Autonomous Coding Agents. *arXiv:2607.12428*, 2026-07-14. <https://arxiv.org/abs/2607.12428>
- [54] Selvanayagam, N., Ghaleb, T. A. AI-to-AI Code Reviews of GitHub Pull Requests. *arXiv:2608.21311*, 2026-08-21. <https://arxiv.org/abs/2608.21311>
- [55] Stolze, M., Strässle, M. When Review Alone No Longer Scales: Layered Supervision in AI-Assisted Software Engineering. *arXiv:2608.26316*, 2026-08-26. <https://arxiv.org/abs/2608.26316>
- [56] Monperrus, M. The End of Code Review: Coding Agents Supersede Human Inspection. *arXiv:2606.13175*, 2026-06-11. <https://arxiv.org/abs/2606.13175>
- [58] Farrag, S. E. The Productivity-Reliability Paradox: Specification-Driven Governance for AI-Augmented Software Development. *arXiv:2605.01160*, 2026-05-01. <https://arxiv.org/abs/2605.01160>