



Contoso Engineering (illustrative mock)

# Implementation guide

From assessment to execution: governance, phased plan, change management and the first 90 days

|                   |            |
|-------------------|------------|
| ASSESSMENT DATE   | 2026-09-27 |
| FRAMEWORK VERSION | 2.0.2      |
| RESPONDENTS       | 14         |
| GENERATED ON      | 2026-09-29 |

# 1. Executive summary



Contoso Engineering (illustrative mock) scored 2.04 (L2 Adopting) from 14 respondents. 9 of 9 dimensions are below the target of 3.0. The plan orders the work by priority (weight × gap) and starts from the lowest-scoring questions of each dimension; each item is done when its L3 anchor is met and confirmed with evidence.

## 1.1 Address before committing to targets

|                                                                                                          |
|----------------------------------------------------------------------------------------------------------|
| <b>Weak foundations:</b> D8                                                                              |
| <b>Perception gap:</b> D1, D2, D3, D5, D6, D7, D8, D9                                                    |
| <b>Low confidence:</b> D6                                                                                |
| <b>Unverified strengths:</b> D1-Q1, D1-Q3, D1-Q4, D1-Q5, D1-Q7, D4-Q1, D4-Q3, D4-Q4, D4-Q5, D4-Q6, D4-Q7 |
| <b>Telemetry below the answers:</b> D4-Q1                                                                |

Wizard inputs: 55% complete.

# 2. Governance

## 2.1 Steering committee

Sponsors the program, approves targets and investment, and removes blockers.

| NAME                | ROLE                                       |
|---------------------|--------------------------------------------|
| Maria Tech Leader   | Champion (maria.tl@cliente-exemplo.com.br) |
| Joao Backend SRE    | Champion (joao.sre@cliente-exemplo.com.br) |
| Sofia Frontend Lead | Champion (sofia.fe@cliente-exemplo.com.br) |

## 2.2 Program office

Runs the plan week to week, keeps the backlog, and reports progress to the steering committee.

To fill with the client

## 2.3 Dimension owners

One accountable owner for each dimension below target.

| ID | DIMENSION                               | PRIORITY  | HORIZON           | OWNER                   |
|----|-----------------------------------------|-----------|-------------------|-------------------------|
| D1 | AI Strategy, Policy and Governance      | P3 Low    | Backlog / monitor | To fill with the client |
| D2 | Enablement, Skills and Culture          | P3 Low    | Backlog / monitor | To fill with the client |
| D3 | Plan, Specify and Design                | P2 Medium | Semester          | To fill with the client |
| D4 | Code and Context Engineering            | P3 Low    | Backlog / monitor | To fill with the client |
| D5 | Review, Quality and Testing             | P2 Medium | Semester          | To fill with the client |
| D6 | Security and AI Supply Chain            | P2 Medium | Semester          | To fill with the client |
| D7 | Deliver and Operate                     | P2 Medium | Semester          | To fill with the client |
| D8 | Engineering Foundations (AI amplifiers) | P1 High   | Next quarter      | To fill with the client |
| D9 | Measurement, Value and AI FinOps        | P2 Medium | Semester          | To fill with the client |

## 2.4 RACI

To fill with the client

## 3. Phased plan

Each dimension below target sits in a horizon set by its priority: P0 in 30 days, P1 next quarter, P2 this semester, P3 in the backlog. Start from the questions listed; an item is done when its L3 anchor holds and the evidence listed confirms it.

### 3.1 P1: Next quarter

#### D8: Engineering Foundations (AI amplifiers) (1.09 → 3.0, P1 High)

| ID    | QUESTION                                                                                        | SCORE | DONE WHEN (L3 ANCHOR)                                                                                           | EVIDENCE TO COLLECT                                            |
|-------|-------------------------------------------------------------------------------------------------|-------|-----------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------|
| D8-Q7 | Automated testing as a control system<br>KPI: % of services meeting the L3 anchor               | 0.75  | Most services have layered automated tests that run on every PR within agreed time budgets. [3]                 | Test suite inventory, pipeline durations, escaped-defect data. |
| D8-Q4 | Healthy data ecosystem<br>KPI: Level of the organization-wide practice                          | 0.92  | Key data is cataloged with owners and quality indicators; most questions can be answered within an hour. [1][2] | Data catalog, quality dashboards.                              |
| D8-Q5 | Reproducible environments for humans and agents<br>KPI: % of repositories meeting the L3 anchor | 1.00  | Most repositories define a reproducible environment; agents use the same definition. [8]                        | Devcontainer files, environment start-up times.                |

#### Strategies for this horizon

- **S2 Foundry + SRE, S3 App Modernization, S1 GitHub Migration:** Start with D8-Q7 (Automated testing as a control system). Target the L3 anchor: Most services have layered automated tests that run on every PR within agreed time budgets.

## 3.2 P2: Semester

### D3: Plan, Specify and Design (2.09 → 3.0, P2 Medium)

| ID    | QUESTION                                                                     | SCORE | DONE WHEN (L3 ANCHOR)                                                                                                                         | EVIDENCE TO COLLECT                                        |
|-------|------------------------------------------------------------------------------|-------|-----------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------|
| D3-Q1 | AI in backlog refinement<br>KPI: % of teams meeting the L3 anchor            | 2.00  | Most teams use AI to draft or improve work items; acceptance criteria are mandatory before work starts. [16][17]                              | Issue templates, sample work items, quality checks.        |
| D3-Q2 | Specification before implementation<br>KPI: % of teams meeting the L3 anchor | 2.00  | Plans or specs are stored in the repository or linked to the issue, and are reviewed by a human before agent implementation. [23][24][27][58] | Spec files, plan reviews, PRs that reference specs.        |
| D3-Q5 | User-centric focus<br>KPI: % of teams meeting the L3 anchor                  | 2.07  | Work items reference the user problem and success measure; feedback is reviewed before prioritization. [1][2][3]                              | Product briefs, feedback loop records, outcome dashboards. |

### D5: Review, Quality and Testing (2.10 → 3.0, P2 Medium)

| ID    | QUESTION                                                                                 | SCORE | DONE WHEN (L3 ANCHOR)                                                                                                                            | EVIDENCE TO COLLECT                                                     |
|-------|------------------------------------------------------------------------------------------|-------|--------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------|
| D5-Q6 | Verification culture and calibrated trust<br>KPI: % of teams meeting the L3 anchor       | 1.77  | Review guidelines explain what to check in AI output; trust in AI output is part of the developer survey. [3][23][35][37][48][52]                | Review guidelines, survey results, defect analysis.                     |
| D5-Q3 | Same quality gates for AI and human code<br>KPI: % of repositories meeting the L3 anchor | 2.00  | Required checks are enforced by rulesets on the protected branches of most repositories, with no bypass for agent identities. [3][7][31][50][55] | Rulesets, required checks, bypass lists.                                |
| D5-Q1 | AI-assisted code review<br>KPI: % of repositories meeting the L3 anchor                  | 2.07  | AI review runs automatically on most PRs; teams track useful vs dismissed suggestions. [6][12][14][54][55]                                       | Repository rulesets, code review adoption metrics, suggestion outcomes. |

## D6: Security and AI Supply Chain (1.90 → 3.0, P2 Medium)

| ID    | QUESTION                                                                               | SCORE | DONE WHEN (L3 ANCHOR)                                                                                                                    | EVIDENCE TO COLLECT                                            |
|-------|----------------------------------------------------------------------------------------|-------|------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------|
| D6-Q5 | AI supply chain<br>KPI: % of repositories meeting the L3 anchor                        | 1.50  | A review process covers AI components; SBOMs and build provenance are produced for most builds. [38][40][43][52]                         | Component review records, SBOM and attestation samples.        |
| D6-Q6 | Threat modeling for AI features and agents<br>KPI: % of services meeting the L3 anchor | 1.62  | Threat models are required for new AI features and agent workflows and reviewed by security. [38][39][40]                                | Threat model documents, security review records.               |
| D6-Q4 | Least privilege for agents<br>KPI: Level of the organization-wide practice             | 1.71  | Agent permissions are documented and reviewed; agents cannot reach production credentials or push to protected branches. [7][19][38][39] | Agent environment configuration, token scopes, access reviews. |

## D7: Deliver and Operate (1.75 → 3.0, P2 Medium)

| ID    | QUESTION                                                                  | SCORE | DONE WHEN (L3 ANCHOR)                                                                                                         | EVIDENCE TO COLLECT                                                       |
|-------|---------------------------------------------------------------------------|-------|-------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------|
| D7-Q4 | Observability of agents<br>KPI: Level of the organization-wide practice   | 1.57  | Agent runs emit telemetry to the central observability stack; dashboards show failures and cost per agent. [19][32][45]       | Telemetry dashboards, alert rules.                                        |
| D7-Q1 | AI in CI/CD pipelines<br>KPI: % of repositories meeting the L3 anchor     | 1.69  | Pipelines are code in most repositories; AI-assisted failure analysis is available to all teams. [16][17]                     | Pipeline repositories, failure-analysis usage, build metrics.             |
| D7-Q3 | AI-assisted incident response<br>KPI: % of services meeting the L3 anchor | 1.71  | On-call engineers in most teams use AI for triage and summaries; post-incident reviews record whether AI helped. [16][17][22] | Incident tooling configuration, incident timelines, time-to-restore data. |

## D9: Measurement, Value and AI FinOps (1.61 → 3.0, P2 Medium)

| ID    | QUESTION                                                                          | SCORE | DONE WHEN (L3 ANCHOR)                                                                                                                              | EVIDENCE TO COLLECT                                             |
|-------|-----------------------------------------------------------------------------------|-------|----------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------|
| D9-Q3 | Pull request flow metrics<br>KPI: Level of the organization-wide practice         | 1.33  | PR lifecycle metrics are reported per organization; agent-authored PRs are identified separately, including their post-merge fix rate. [6][34][50] | PR lifecycle metrics, agent PR reports, follow-up fix analysis. |
| D9-Q7 | Business value linkage<br>KPI: Level of the organization-wide practice            | 1.54  | A business case with explicit assumptions exists and is reviewed at least yearly. [18][28][32]                                                     | Business case, value reports.                                   |
| D9-Q4 | Developer experience and friction<br>KPI: Level of the organization-wide practice | 1.57  | A survey runs at least twice a year with good participation; results are shared and acted on. [2][44]                                              | Survey instrument, participation rate, action log.              |

### Strategies for this horizon

- **S5 GitHub Copilot Acceleration:** Start with D9-Q3 (Pull request flow metrics). Target the L3 anchor: PR lifecycle metrics are reported per organization; agent-authored PRs are identified separately, including their post-merge fix rate.
- **S6 Agentic Activation, S7 Security & Governance:** Start with D6-Q5 (AI supply chain). Target the L3 anchor: A review process covers AI components; SBOMs and build provenance are produced for most builds.

### 3.3 P3: Backlog / monitor

#### D1: AI Strategy, Policy and Governance (2.82 → 3.0, P3 Low)

| ID    | QUESTION                                                                                         | SCORE | DONE WHEN (L3 ANCHOR)                                                                                                                                                                  | EVIDENCE TO COLLECT                                              |
|-------|--------------------------------------------------------------------------------------------------|-------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------|
| D1-Q5 | Autonomy levels for AI work<br><b>unverified</b><br>KPI: Level of the organization-wide practice | 2.69  | A published matrix maps task types (for example dependency upgrades, test generation, feature work, production changes) to autonomy levels and required approvals. [7][26][32][49][56] | Autonomy matrix, repository rulesets, change records.            |
| D1-Q6 | Responsible AI and risk framework<br>KPI: Level of the organization-wide practice                | 2.71  | A named framework is adopted; AI-related risks are in the risk register with owners and reviews. [20][21][41][42]                                                                      | Framework mapping, risk register entries, audit reports.         |
| D1-Q3 | Approved tools and models <b>unverified</b><br>KPI: Level of the organization-wide practice      | 2.75  | Enterprise/organization policies enable only approved features and models; the catalog lists owner, data handling and review date for each tool. [2][15][32][49]                       | Copilot policy settings, tool catalog, model evaluation records. |

#### D2: Enablement, Skills and Culture (2.28 → 3.0, P3 Low)

| ID    | QUESTION                                                                | SCORE | DONE WHEN (L3 ANCHOR)                                                                                                                                                                                                                                                                  | EVIDENCE TO COLLECT                                              |
|-------|-------------------------------------------------------------------------|-------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------|
| D2-Q1 | Structured AI training<br>KPI: % of engineers meeting the L3 anchor     | 2.15  | Role-based curriculum (developer, reviewer, platform, security) with completion tracked; training is required before agent features are enabled; it teaches learning-preserving patterns (ask for explanations, attempt first, then compare) and not only full delegation. [2][30][48] | Learning paths, completion rates, enablement gating rules.       |
| D2-Q6 | AI-assisted onboarding<br>KPI: % of engineers meeting the L3 anchor     | 2.15  | Onboarding includes AI-guided codebase tours and repository instructions; time to first merged PR is tracked; new engineers are checked on code reading and debugging, not only output. [27][48]                                                                                       | Onboarding playbook, time-to-first-PR data, skill check results. |
| D2-Q4 | Context-engineering skills<br>KPI: % of engineers meeting the L3 anchor | 2.21  | Guidance and examples on context engineering are part of training; teams review their instructions and prompts for quality. [25][30][32]                                                                                                                                               | Guidance pages, review checklists, before/after metrics.         |

## D4: Code and Context Engineering (2.72 → 3.0, P3 Low)

| ID    | QUESTION                                                                                          | SCORE | DONE WHEN (L3 ANCHOR)                                                                                                                                                                                                                         | EVIDENCE TO COLLECT                                                                              |
|-------|---------------------------------------------------------------------------------------------------|-------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------|
| D4-Q4 | Repository instructions<br><b>unverified</b><br>KPI: % of repositories meeting the L3 anchor      | 2.50  | Most active repositories have structured instructions (build, test, conventions, constraints) with an owner; changes go through PR review; files are updated when the codebase changes rather than committed once. [2][8][25][27][46][47][55] | Instruction files, coverage across repositories, change history, before/after agent PR metrics.  |
| D4-Q1 | Depth of AI use across surfaces<br><b>unverified</b><br>KPI: % of engineers meeting the L3 anchor | 2.62  | Many engineers regularly use at least one GitHub agent surface ("Agent first"), confirmed by usage metrics. [6]                                                                                                                               | Copilot usage metrics dashboard or API: adoption cohort distribution, daily/weekly active users. |
| D4-Q2 | Agent mode for multi-file work<br>KPI: % of engineers meeting the L3 anchor                       | 2.64  | Agent mode is the default for multi-file refactors and features in most teams; changes are reviewed in the diff before commit. [6][27]                                                                                                        | Usage by feature/mode, team guidelines.                                                          |

## 4. Change management

### 4.1 ADKAR plan

---

Awareness, Desire, Knowledge, Ability and Reinforcement: the interventions planned for each stage.

ADKAR Change Plan derived from the Learning Survey:

Awareness: Share the consolidated training plan in an all-hands; each dev receives a personalized plan by email.

Desire: Make it visible that workshops have pre-validated attendees (not a generic opt-in).

Knowledge: Top 5 workshops (from the Learning Survey):

1. Copilot Spaces (shared context) (5 attendees)
2. A2A (Agent-to-Agent) protocol (4 attendees)
3. Microsoft Foundry overview (4 attendees)
4. Automatic documentation (4 attendees)
5. Ask mode (effective questions) (3 attendees)

Ability: Biweekly office hours (no fixed agenda; devs bring practical questions).

Reinforcement: Monthly adoption metrics published; recognition for Champions; quarterly plan review with a new Learning Survey.

### 4.2 Communication plan

---

| AUDIENCE                            | CHANNEL     | FREQUENCY        | OWNER             |
|-------------------------------------|-------------|------------------|-------------------|
| Champions Kickoff (3)               | 2h live     | W1               | Eng Lead          |
| Copilot Spaces (shared context) (5) | 4h hands-on | W2               | Maria Tech Leader |
| A2A (Agent-to-Agent) protocol (4)   | 4h hands-on | W3               | Maria Tech Leader |
| Microsoft Foundry overview (4)      | 4h hands-on | W4               | Maria Tech Leader |
| Office hours #1 (Everyone)          | 1h Q&A      | W4               | Champions         |
| Biweekly office hours (Everyone)    | 1h Q&A      | W6, W8, W10, W12 | Champions         |

### 4.3 Enablement and training plan

---

D2 Enablement, Skills and Culture scored 2.28 (L2 Adopting). Build the training cohorts from the Learning and Growth Survey (/training-plan) and track them in D2. Training plan found: output/training-plan-2026-09-29.md.

| AUDIENCE                               | FORMAT                                 | CADENCE                                                 |
|----------------------------------------|----------------------------------------|---------------------------------------------------------|
| DS-D2 (Copilot Adoption): 3 devs       | Self-paced (my pace, more flexibility) | 4-6 week cohort with live sessions and a self-paced lab |
| DS-D3 (MS/GH Tooling Breadth): 3 devs  | Self-paced (my pace, more flexibility) | 4-6 week cohort with live sessions and a self-paced lab |
| DS-D4 (AI Dev Practices): 1 devs       | Self-paced (my pace, more flexibility) | 4-6 week cohort with live sessions and a self-paced lab |
| DS-D5 (Agent Concepts Mastery): 2 devs | Self-paced (my pace, more flexibility) | 4-6 week cohort with live sessions and a self-paced lab |
| DS-D6 (Instructions Maturity): 4 devs  | Self-paced (my pace, more flexibility) | 4-6 week cohort with live sessions and a self-paced lab |
| DS-D7 (Best Practices): 3 devs         | Self-paced (my pace, more flexibility) | 4-6 week cohort with live sessions and a self-paced lab |
| DS-D8 (Security & Governance): 2 devs  | Self-paced (my pace, more flexibility) | 4-6 week cohort with live sessions and a self-paced lab |

# 5. Risks

Risks raised by the scoring rules, then the client's own risk register.

| RISK                                                                                                                                                               | AFFECTED                                                                    | MITIGATION                                                                                      |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------|
| <b>Weak foundations</b><br>Amplifier dimensions are at least one band below the overall level. DORA finds that AI amplifies existing strengths and weaknesses [3]. | D8                                                                          | Raise these dimensions before scaling AI adoption.                                              |
| <b>Perception gap</b><br>Executives rate these dimensions at least one band above hands-on engineers.                                                              | D1, D2, D3, D5, D6, D7, D8, D9                                              | Review the gap with both groups and confirm with telemetry and evidence before setting targets. |
| <b>Low confidence</b><br>More than 30% of the answers are NA in these dimensions.                                                                                  | D6                                                                          | Find respondents who know these practices and ask again.                                        |
| <b>Unverified strengths</b><br>Scored L3 or higher, but fewer than half of the answers carry evidence.                                                             | D1-Q1, D1-Q3, D1-Q4, D1-Q5, D1-Q7, D4-Q1, D4-Q3, D4-Q4, D4-Q5, D4-Q6, D4-Q7 | Confirm with evidence before presenting these questions as strengths.                           |
| <b>Telemetry below the answers</b><br>Copilot usage metrics support a lower D4-Q1 level than the answers.                                                          | D4-Q1                                                                       | Review D4-Q1 with the telemetry and set adoption targets from it.                               |

## 5.1 Client risk register

To fill with the client

## 6. Success metrics and re-assessment

Current score and target for each dimension below target. The KPI of each item is in the phased plan.

| ID | DIMENSION                               | SCORE | TARGET | GAP  | HORIZON           |
|----|-----------------------------------------|-------|--------|------|-------------------|
| D1 | AI Strategy, Policy and Governance      | 2.82  | 3.0    | 0.18 | Backlog / monitor |
| D2 | Enablement, Skills and Culture          | 2.28  | 3.0    | 0.72 | Backlog / monitor |
| D3 | Plan, Specify and Design                | 2.09  | 3.0    | 0.91 | Semester          |
| D4 | Code and Context Engineering            | 2.72  | 3.0    | 0.28 | Backlog / monitor |
| D5 | Review, Quality and Testing             | 2.10  | 3.0    | 0.90 | Semester          |
| D6 | Security and AI Supply Chain            | 1.90  | 3.0    | 1.10 | Semester          |
| D7 | Deliver and Operate                     | 1.75  | 3.0    | 1.25 | Semester          |
| D8 | Engineering Foundations (AI amplifiers) | 1.09  | 3.0    | 1.91 | Next quarter      |
| D9 | Measurement, Value and AI FinOps        | 1.61  | 3.0    | 1.39 | Semester          |

- Re-assess at the end of each horizon with the same form, then compare rounds with make compare BEFORE=... AFTER=... (the comparison PDF shows the deltas).
- Where available, confirm D4 and D9 with Copilot usage metrics [6] and delivery metrics (D9-Q2) before reporting progress.

## 7. First 90 days

### 7.1 Weeks 1 to 4

---

1. W1: Champions Kickoff
2. W2: Copilot Spaces (shared context)
3. W3: A2A (Agent-to-Agent) protocol
4. W4: Microsoft Foundry overview
5. W4: Office hours #1

### 7.2 Weeks 5 to 8

---

- Start S2 Foundry + SRE: Start with D8-Q7 (Automated testing as a control system). Target the L3 anchor: Most services have layered automated tests that run on every PR within agreed time budgets.
- Start S5 GitHub Copilot Acceleration: Start with D9-Q3 (Pull request flow metrics). Target the L3 anchor: PR lifecycle metrics are reported per organization; agent-authored PRs are identified separately, including their post-merge fix rate.

### 7.3 Weeks 9 to 12

---

1. W10: Biweekly office hours

## 8. References

- [1] DORA. DORA AI Capabilities Model. Google Cloud, 2025. <https://dora.dev/ai/capabilities-model/report/>
- [2] DORA. DORA AI Capabilities Model: Survey Questions. Google Cloud, 2025. <https://dora.dev/ai/capabilities-model/questions/>
- [3] Google Cloud. Announcing the 2025 DORA Report: State of AI-Assisted Software Development. 2025. <https://cloud.google.com/blog/products/ai-machine-learning/announcing-the-2025-dora-report>
- [6] GitHub Docs. GitHub Copilot usage metrics. <https://docs.github.com/en/copilot/concepts/billing-and-usage/copilot-usage-metrics/copilot-metrics>
- [7] GitHub Docs. Risks and mitigations for GitHub Copilot cloud agent. <https://docs.github.com/en/enterprise-cloud@latest/copilot/concepts/security-governance-and-network-settings/risks-and-mitigations>
- [8] GitHub Docs. Best practices for using GitHub Copilot to work on tasks. <https://docs.github.com/en/enterprise-cloud@latest/copilot/tutorials/cloud-agent/get-the-best-results>
- [12] GitHub Docs. About GitHub Copilot code review. <https://docs.github.com/en/copilot/concepts/agents/code-review>
- [14] GitHub Docs. Application card: GitHub Copilot Agents. <https://docs.github.com/en/copilot/responsible-use/agents>
- [15] GitHub Docs. Managing policies and features for GitHub Copilot in your organization. <https://docs.github.com/en/copilot/how-tos/administer-copilot/manage-for-organization/manage-policies>
- [16] Microsoft Azure Blog. Agentic DevOps: Evolving software development with GitHub Copilot and Microsoft Azure. 2025. <https://azure.microsoft.com/en-us/blog/agentic-devops-evolving-software-development-with-github-copilot-and-microsoft-azure/>
- [17] Microsoft for Developers. Agentic DevOps in action: Reimagining every phase of the developer lifecycle. 2025. <https://developer.microsoft.com/blog/reimagining-every-phase-of-the-developer-lifecycle/>
- [18] Microsoft Learn. AI strategy: Cloud Adoption Framework. <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/ai/strategy>
- [19] Microsoft Learn. Govern and secure AI agents: Cloud Adoption Framework. <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/ai-agents/governance-security-across-organization>
- [20] Microsoft Learn. Responsible AI policies: Cloud Adoption Framework. <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/ai/responsible-ai-policies>
- [21] Microsoft. Responsible AI Principles and Approach. <https://www.microsoft.com/en-us/ai/principles-and-approach>
- [22] Microsoft Azure Blog. Announcing Azure Copilot agents and AI infrastructure innovations. 2025. <https://azure.microsoft.com/en-us/blog/announcing-azure-copilot-agents-and-ai-infrastructure-innovations/>
- [23] Anthropic. Anthropic Economic Index: AI's impact on software development. 2025-04-28. <https://www.anthropic.com/research/impact-software-development>
- [24] Anthropic. How Claude Code is used in practice. 2026-06-16. <https://www.anthropic.com/research/claude-code-expertise>
- [25] Anthropic. Effective context engineering for AI agents. 2025-09-29. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>
- [26] Anthropic. Building effective agents. 2024-12-19. <https://www.anthropic.com/engineering/building-effective-agents>
- [27] Anthropic. Best practices for Claude Code. Claude Code Docs. <https://code.claude.com/docs/en/best-practices>
- [28] Gartner. Gartner Says 75% of Enterprise Software Engineers Will Use AI Code Assistants by 2028. 2024-04-11. <https://www.gartner.com/en/newsroom/press-releases/2024-04-11-gartner-says-75-percent-of-enterprise-software-engineers-will-use-ai-code-assistants-by-2028>
- [30] Gartner. Gartner Says Generative AI will Require 80% of Engineering Workforce to Upskill Through 2027. 2024-10-03. <https://www.gartner.com/en/newsroom/press-releases/2024-10-03-gartner-says-generative-ai-will-require-80-percent-of-engineering-workforce-to-upskill-through-2027>

- [31] Gartner. Gartner Says the Market for Enterprise AI Coding Agents Is Entering a New Phase of Expansion and Competitive Realignment. 2026-05-20. <https://www.gartner.com/en/newsroom/press-releases/2026-05-20-gartner-says-the-market-for-enterprise-ai-coding-agents-is-entering-a-new-phase-of-expansion-and-competitive-realignment>
- [32] Gartner. Gartner Predicts AI Coding Costs Will Surpass Average Developer's Salary by 2028 as Token Consumption Surges. 2026-06-24. <https://www.gartner.com/en/newsroom/press-releases/2026-06-24-gartner-predicts-ai-coding-costs-will-surpass-average-developer-salary-by-2028-as-token-consumption-surges>
- [34] Cui, Z. K., Demirer, M., Jaffe, S., Musolff, L., Peng, S., Salz, T. The Effects of Generative AI on High-Skilled Work: Evidence from Three Field Experiments with Software Developers. Management Science, published online 2026-02-27. <https://doi.org/10.1287/mnsc.2025.00535> (pre-print: [https://papers.ssrn.com/sol3/papers.cfm?abstract\\_id=4945566](https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4945566))
- [35] METR. Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity. 2025-07-10. <https://metr.org/blog/2025-07-10-early-2025-ai-experienced-os-dev-study/> (paper: <https://arxiv.org/abs/2507.09089>)
- [37] Stack Overflow. 2025 Developer Survey: AI. <https://survey.stackoverflow.co/2025/ai>
- [38] OWASP GenAI Security Project. OWASP Top 10 for LLM Applications 2025. <https://genai.owasp.org/resource/owasp-top-10-for-llm-applications-2025/>
- [39] OWASP GenAI Security Project. OWASP Top 10 for Agentic Applications for 2026. 2025-12-09. <https://genai.owasp.org/resource/owasp-top-10-for-agentic-applications-for-2026/>
- [40] NIST. SP 800-218A: Secure Software Development Practices for Generative AI and Dual-Use Foundation Models: An SSDF Community Profile. 2024-07. <https://csrc.nist.gov/pubs/sp/800/218/a/final>
- [41] NIST. AI Risk Management Framework. <https://www.nist.gov/itl/ai-risk-management-framework>
- [42] ISO. ISO/IEC 42001:2023: AI management systems. <https://www.iso.org/standard/42001>
- [43] OpenSSF. SLSA: Supply-chain Levels for Software Artifacts. <https://slsa.dev/>
- [44] Forsgren, N., Storey, M.-A., Maddila, C., Zimmermann, T., Houck, B., Butler, J. The SPACE of Developer Productivity. ACM Queue 19(1), 2021-03-06. <https://queue.acm.org/detail.cfm?id=3454124> (DOI: <https://doi.org/10.1145/3454122.3454124>)
- [45] Liu, B., Qiu, H., Goiri, Í., Fonseca, R., Bianchini, R., Choukse, E. Agentic Coding in the Wild: Characterizing GitHub Copilot Traces at Production Scale. arXiv:2608.00101, 2026-07-30. <https://arxiv.org/abs/2608.00101>
- [46] Arabat, A., Sayagh, M. Toward Instructions-as-Code: Understanding the Impact of Instruction Files on Agentic Pull Requests. arXiv:2606.13449, 2026-06-11. <https://arxiv.org/abs/2606.13449>
- [47] Denisov-Blanch, Y., Agarwal, S., Azaletskiy, P., He, H., Schaeffer, R., Miranda, B., Vasilescu, B., Koyejo, S. A Few Pages of Markdown: Committed AI Configuration and Lower Quality Cost after Coding-Agent Adoption. arXiv:2608.25241, 2026-08-26. <https://arxiv.org/abs/2608.25241>
- [48] Shen, J. H., Tamkin, A. How AI Impacts Skill Formation. arXiv:2601.20245, 2026-01-28. <https://arxiv.org/abs/2601.20245> (Anthropic summary: <https://www.anthropic.com/research/AI-assistance-coding-skills>)
- [49] Pinna, G., Gong, J., Williams, D., Sarro, F. Comparing AI Coding Agents: A Task-Stratified Analysis of Pull Request Acceptance. arXiv:2602.08915, 2026-02-09. <https://arxiv.org/abs/2602.08915>
- [50] Takerngsaksiri, W., Duong, N., Barnett, S. Who Finishes the Job? A Study of Follow-Up Fixes and Commit Authorship on AI Coding Agent Pull Requests. arXiv:2609.26847, 2026-09-22. <https://arxiv.org/abs/2609.26847>
- [52] Sakib, A. H. M. N., Banik, D., Jadliwala, M. Trust but Verify? Uncovering the Security Debt of Autonomous Coding Agents. arXiv:2607.12428, 2026-07-14. <https://arxiv.org/abs/2607.12428>
- [54] Selvanayagam, N., Ghaleb, T. A. AI-to-AI Code Reviews of GitHub Pull Requests. arXiv:2608.21311, 2026-08-21. <https://arxiv.org/abs/2608.21311>
- [55] Stolze, M., Strässle, M. When Review Alone No Longer Scales: Layered Supervision in AI-Assisted Software Engineering. arXiv:2608.26316, 2026-08-26. <https://arxiv.org/abs/2608.26316>
- [56] Monperrus, M. The End of Code Review: Coding Agents Supersede Human Inspection. arXiv:2606.13175, 2026-06-11. <https://arxiv.org/abs/2606.13175>
- [58] Farrag, S. E. The Productivity-Reliability Paradox: Specification-Driven Governance for AI-Augmented Software Development. arXiv:2605.01160, 2026-05-01. <https://arxiv.org/abs/2605.01160>