



Contoso Engineering (illustrative mock)

Implementation guide

From assessment to execution: governance, phased plan, change management and the first 90 days

ASSESSMENT DATE	2026-09-27
FRAMEWORK VERSION	2.0.1
RESPONDENTS	14
GENERATED ON	2026-09-29

1. Executive summary



Contoso Engineering (illustrative mock) scored 2.04 (L2 Adopting) from 14 respondents. 9 of 9 dimensions are below the target of 3.0. The plan orders the work by priority (weight × gap) and starts from the lowest-scoring questions of each dimension; each item is done when its L3 anchor is met and confirmed with evidence.

1.1 Address before committing to targets

Weak foundations: D8
Perception gap: D1, D2, D3, D5, D6, D7, D8, D9
Low confidence: D6
Unverified strengths: D1-Q1, D1-Q3, D1-Q4, D1-Q5, D1-Q7, D4-Q1, D4-Q3, D4-Q4, D4-Q5, D4-Q6, D4-Q7
Telemetry below the answers: D4-Q1

Wizard inputs: 55% complete.

2. Governance

2.1 Steering committee

Sponsors the program, approves targets and investment, and removes blockers.

NAME	ROLE
Maria Tech Leader	Champion (maria.tl@cliente-exemplo.com.br)
Joao Backend SRE	Champion (joao.sre@cliente-exemplo.com.br)
Sofia Frontend Lead	Champion (sofia.fe@cliente-exemplo.com.br)

2.2 Program office

Runs the plan week to week, keeps the backlog, and reports progress to the steering committee.

To fill with the client

2.3 Dimension owners

One accountable owner for each dimension below target.

ID	DIMENSION	PRIORITY	HORIZON	OWNER
D1	AI Strategy, Policy and Governance	P3 Low	Backlog / monitor	To fill with the client
D2	Enablement, Skills and Culture	P3 Low	Backlog / monitor	To fill with the client
D3	Plan, Specify and Design	P2 Medium	Semester	To fill with the client
D4	Code and Context Engineering	P3 Low	Backlog / monitor	To fill with the client
D5	Review, Quality and Testing	P2 Medium	Semester	To fill with the client
D6	Security and AI Supply Chain	P2 Medium	Semester	To fill with the client
D7	Deliver and Operate	P2 Medium	Semester	To fill with the client
D8	Engineering Foundations (AI amplifiers)	P1 High	Next quarter	To fill with the client
D9	Measurement, Value and AI FinOps	P2 Medium	Semester	To fill with the client

2.4 RACI

To fill with the client

3. Phased plan

Each dimension below target sits in a horizon set by its priority: P0 in 30 days, P1 next quarter, P2 this semester, P3 in the backlog. Start from the questions listed; an item is done when its L3 anchor holds and the evidence listed confirms it.

3.1 P1: Next quarter

D8: Engineering Foundations (AI amplifiers) (1.09 → 3.0, P1 High)

ID	QUESTION	SCORE	DONE WHEN (L3 ANCHOR)	EVIDENCE TO COLLECT
D8-Q7	Automated testing as a control system KPI: % of services meeting the L3 anchor	0.75	Most services have layered automated tests that run on every PR within agreed time budgets. [3]	Test suite inventory, pipeline durations, escaped-defect data.
D8-Q4	Healthy data ecosystem KPI: Level of the organization-wide practice	0.92	Key data is cataloged with owners and quality indicators; most questions can be answered within an hour. [1][2]	Data catalog, quality dashboards.
D8-Q5	Reproducible environments for humans and agents KPI: % of repositories meeting the L3 anchor	1.00	Most repositories define a reproducible environment; agents use the same definition. [8]	Devcontainer files, environment start-up times.

Strategies for this horizon

- **S2 Foundry + SRE, S3 App Modernization, S1 GitHub Migration:** Start with D8-Q7 (Automated testing as a control system). Target the L3 anchor: Most services have layered automated tests that run on every PR within agreed time budgets.

3.2 P2: Semester

D3: Plan, Specify and Design (2.09 → 3.0, P2 Medium)

ID	QUESTION	SCORE	DONE WHEN (L3 ANCHOR)	EVIDENCE TO COLLECT
D3-Q1	AI in backlog refinement KPI: % of teams meeting the L3 anchor	2.00	Most teams use AI to draft or improve work items; acceptance criteria are mandatory before work starts. [16][17]	Issue templates, sample work items, quality checks.
D3-Q2	Specification before implementation KPI: % of teams meeting the L3 anchor	2.00	Plans or specs are stored in the repository or linked to the issue, and are reviewed by a human before agent implementation. [23][24][27][58]	Spec files, plan reviews, PRs that reference specs.
D3-Q5	User-centric focus KPI: % of teams meeting the L3 anchor	2.07	Work items reference the user problem and success measure; feedback is reviewed before prioritization. [1][2][3]	Product briefs, feedback loop records, outcome dashboards.

D5: Review, Quality and Testing (2.10 → 3.0, P2 Medium)

ID	QUESTION	SCORE	DONE WHEN (L3 ANCHOR)	EVIDENCE TO COLLECT
D5-Q6	Verification culture and calibrated trust KPI: % of teams meeting the L3 anchor	1.77	Review guidelines explain what to check in AI output; trust in AI output is part of the developer survey. [3][23][35][37][48][52]	Review guidelines, survey results, defect analysis.
D5-Q3	Same quality gates for AI and human code KPI: % of repositories meeting the L3 anchor	2.00	Required checks are enforced by rulesets on the protected branches of most repositories, with no bypass for agent identities. [3][7][31][50][55]	Rulesets, required checks, bypass lists.
D5-Q1	AI-assisted code review KPI: % of repositories meeting the L3 anchor	2.07	AI review runs automatically on most PRs; teams track useful vs dismissed suggestions. [6][12][14][54][55]	Repository rulesets, code review adoption metrics, suggestion outcomes.

D6: Security and AI Supply Chain (1.90 → 3.0, P2 Medium)

ID	QUESTION	SCORE	DONE WHEN (L3 ANCHOR)	EVIDENCE TO COLLECT
D6-Q5	AI supply chain KPI: % of repositories meeting the L3 anchor	1.50	A review process covers AI components; SBOMs and build provenance are produced for most builds. [38][40][43][52]	Component review records, SBOM and attestation samples.
D6-Q6	Threat modeling for AI features and agents KPI: % of services meeting the L3 anchor	1.62	Threat models are required for new AI features and agent workflows and reviewed by security. [38][39][40]	Threat model documents, security review records.
D6-Q4	Least privilege for agents KPI: Level of the organization-wide practice	1.71	Agent permissions are documented and reviewed; agents cannot reach production credentials or push to protected branches. [7][19][38][39]	Agent environment configuration, token scopes, access reviews.

D7: Deliver and Operate (1.75 → 3.0, P2 Medium)

ID	QUESTION	SCORE	DONE WHEN (L3 ANCHOR)	EVIDENCE TO COLLECT
D7-Q4	Observability of agents KPI: Level of the organization-wide practice	1.57	Agent runs emit telemetry to the central observability stack; dashboards show failures and cost per agent. [19][32][45]	Telemetry dashboards, alert rules.
D7-Q1	AI in CI/CD pipelines KPI: % of repositories meeting the L3 anchor	1.69	Pipelines are code in most repositories; AI-assisted failure analysis is available to all teams. [16][17]	Pipeline repositories, failure-analysis usage, build metrics.
D7-Q3	AI-assisted incident response KPI: % of services meeting the L3 anchor	1.71	On-call engineers in most teams use AI for triage and summaries; post-incident reviews record whether AI helped. [16][17][22]	Incident tooling configuration, incident timelines, time-to-restore data.

D9: Measurement, Value and AI FinOps (1.61 → 3.0, P2 Medium)

ID	QUESTION	SCORE	DONE WHEN (L3 ANCHOR)	EVIDENCE TO COLLECT
D9-Q3	Pull request flow metrics KPI: Level of the organization-wide practice	1.33	PR lifecycle metrics are reported per organization; agent-authored PRs are identified separately, including their post-merge fix rate. [6][34][50]	PR lifecycle metrics, agent PR reports, follow-up fix analysis.
D9-Q7	Business value linkage KPI: Level of the organization-wide practice	1.54	A business case with explicit assumptions exists and is reviewed at least yearly. [18][28][32]	Business case, value reports.
D9-Q4	Developer experience and friction KPI: Level of the organization-wide practice	1.57	A survey runs at least twice a year with good participation; results are shared and acted on. [2][44]	Survey instrument, participation rate, action log.

Strategies for this horizon

- **S5 GitHub Copilot Acceleration:** Start with D9-Q3 (Pull request flow metrics). Target the L3 anchor: PR lifecycle metrics are reported per organization; agent-authored PRs are identified separately, including their post-merge fix rate.
- **S6 Agentic Activation, S7 Security & Governance:** Start with D6-Q5 (AI supply chain). Target the L3 anchor: A review process covers AI components; SBOMs and build provenance are produced for most builds.

3.3 P3: Backlog / monitor

D1: AI Strategy, Policy and Governance (2.82 → 3.0, P3 Low)

ID	QUESTION	SCORE	DONE WHEN (L3 ANCHOR)	EVIDENCE TO COLLECT
D1-Q5	Autonomy levels for AI work unverified KPI: Level of the organization-wide practice	2.69	A published matrix maps task types (for example dependency upgrades, test generation, feature work, production changes) to autonomy levels and required approvals. [7][26][32][49][56]	Autonomy matrix, repository rulesets, change records.
D1-Q6	Responsible AI and risk framework KPI: Level of the organization-wide practice	2.71	A named framework is adopted; AI-related risks are in the risk register with owners and reviews. [20][21][41][42]	Framework mapping, risk register entries, audit reports.
D1-Q3	Approved tools and models unverified KPI: Level of the organization-wide practice	2.75	Enterprise/organization policies enable only approved features and models; the catalog lists owner, data handling and review date for each tool. [2][15][32][49]	Copilot policy settings, tool catalog, model evaluation records.

D2: Enablement, Skills and Culture (2.28 → 3.0, P3 Low)

ID	QUESTION	SCORE	DONE WHEN (L3 ANCHOR)	EVIDENCE TO COLLECT
D2-Q1	Structured AI training KPI: % of engineers meeting the L3 anchor	2.15	Role-based curriculum (developer, reviewer, platform, security) with completion tracked; training is required before agent features are enabled; it teaches learning-preserving patterns (ask for explanations, attempt first, then compare) and not only full delegation. [2][30][48]	Learning paths, completion rates, enablement gating rules.
D2-Q6	AI-assisted onboarding KPI: % of engineers meeting the L3 anchor	2.15	Onboarding includes AI-guided codebase tours and repository instructions; time to first merged PR is tracked; new engineers are checked on code reading and debugging, not only output. [27][48]	Onboarding playbook, time-to-first-PR data, skill check results.
D2-Q4	Context-engineering skills KPI: % of engineers meeting the L3 anchor	2.21	Guidance and examples on context engineering are part of training; teams review their instructions and prompts for quality. [25][30][32]	Guidance pages, review checklists, before/after metrics.

D4: Code and Context Engineering (2.72 → 3.0, P3 Low)

ID	QUESTION	SCORE	DONE WHEN (L3 ANCHOR)	EVIDENCE TO COLLECT
D4-Q4	Repository instructions unverified KPI: % of repositories meeting the L3 anchor	2.50	Most active repositories have structured instructions (build, test, conventions, constraints) with an owner; changes go through PR review; files are updated when the codebase changes rather than committed once. [2][8][25][27][46][47][55]	Instruction files, coverage across repositories, change history, before/after agent PR metrics.
D4-Q1	Depth of AI use across surfaces unverified KPI: % of engineers meeting the L3 anchor	2.62	Many engineers regularly use at least one GitHub agent surface ("Agent first"), confirmed by usage metrics. [6]	Copilot usage metrics dashboard or API: adoption cohort distribution, daily/weekly active users.
D4-Q2	Agent mode for multi-file work KPI: % of engineers meeting the L3 anchor	2.64	Agent mode is the default for multi-file refactors and features in most teams; changes are reviewed in the diff before commit. [6][27]	Usage by feature/mode, team guidelines.

4. Change management

4.1 ADKAR plan

Awareness, Desire, Knowledge, Ability and Reinforcement: the interventions planned for each stage.

ADKAR Change Plan derived from the Learning Survey:

Awareness: Share the consolidated training plan in an all-hands; each dev receives a personalized plan by email.

Desire: Make it visible that workshops have pre-validated attendees (not a generic opt-in).

Knowledge: Top 5 workshops (from the Learning Survey):

1. Copilot Spaces (shared context) (5 attendees)
2. A2A (Agent-to-Agent) protocol (4 attendees)
3. Microsoft Foundry overview (4 attendees)
4. Automatic documentation (4 attendees)
5. Ask mode (effective questions) (3 attendees)

Ability: Biweekly office hours (no fixed agenda; devs bring practical questions).

Reinforcement: Monthly adoption metrics published; recognition for Champions; quarterly plan review with a new Learning Survey.

4.2 Communication plan

AUDIENCE	CHANNEL	FREQUENCY	OWNER
Champions Kickoff (3)	2h live	W1	Eng Lead
Copilot Spaces (shared context) (5)	4h hands-on	W2	Maria Tech Leader
A2A (Agent-to-Agent) protocol (4)	4h hands-on	W3	Maria Tech Leader
Microsoft Foundry overview (4)	4h hands-on	W4	Maria Tech Leader
Office hours #1 (Everyone)	1h Q&A	W4	Champions
Biweekly office hours (Everyone)	1h Q&A	W6, W8, W10, W12	Champions

4.3 Enablement and training plan

D2 Enablement, Skills and Culture scored 2.28 (L2 Adopting). Build the training cohorts from the Learning and Growth Survey (/training-plan) and track them in D2. Training plan found: output/training-plan-2026-09-29.md.

AUDIENCE	FORMAT	CADENCE
DS-D2 (Copilot Adoption): 3 devs	Self-paced (my pace, more flexibility)	4-6 week cohort with live sessions and a self-paced lab
DS-D3 (MS/GH Tooling Breadth): 3 devs	Self-paced (my pace, more flexibility)	4-6 week cohort with live sessions and a self-paced lab
DS-D4 (AI Dev Practices): 1 devs	Self-paced (my pace, more flexibility)	4-6 week cohort with live sessions and a self-paced lab
DS-D5 (Agent Concepts Mastery): 2 devs	Self-paced (my pace, more flexibility)	4-6 week cohort with live sessions and a self-paced lab
DS-D6 (Instructions Maturity): 4 devs	Self-paced (my pace, more flexibility)	4-6 week cohort with live sessions and a self-paced lab
DS-D7 (Best Practices): 3 devs	Self-paced (my pace, more flexibility)	4-6 week cohort with live sessions and a self-paced lab
DS-D8 (Security & Governance): 2 devs	Self-paced (my pace, more flexibility)	4-6 week cohort with live sessions and a self-paced lab

5. Risks

Risks raised by the scoring rules, then the client's own risk register.

RISK	AFFECTED	MITIGATION
Weak foundations Amplifier dimensions are at least one band below the overall level. DORA finds that AI amplifies existing strengths and weaknesses [3].	D8	Raise these dimensions before scaling AI adoption.
Perception gap Executives rate these dimensions at least one band above hands-on engineers.	D1, D2, D3, D5, D6, D7, D8, D9	Review the gap with both groups and confirm with telemetry and evidence before setting targets.
Low confidence More than 30% of the answers are NA in these dimensions.	D6	Find respondents who know these practices and ask again.
Unverified strengths Scored L3 or higher, but fewer than half of the answers carry evidence.	D1-Q1, D1-Q3, D1-Q4, D1-Q5, D1-Q7, D4-Q1, D4-Q3, D4-Q4, D4-Q5, D4-Q6, D4-Q7	Confirm with evidence before presenting these questions as strengths.
Telemetry below the answers Copilot usage metrics support a lower D4-Q1 level than the answers.	D4-Q1	Review D4-Q1 with the telemetry and set adoption targets from it.

5.1 Client risk register

To fill with the client

6. Success metrics and re-assessment

Current score and target for each dimension below target. The KPI of each item is in the phased plan.

ID	DIMENSION	SCORE	TARGET	GAP	HORIZON
D1	AI Strategy, Policy and Governance	2.82	3.0	0.18	Backlog / monitor
D2	Enablement, Skills and Culture	2.28	3.0	0.72	Backlog / monitor
D3	Plan, Specify and Design	2.09	3.0	0.91	Semester
D4	Code and Context Engineering	2.72	3.0	0.28	Backlog / monitor
D5	Review, Quality and Testing	2.10	3.0	0.90	Semester
D6	Security and AI Supply Chain	1.90	3.0	1.10	Semester
D7	Deliver and Operate	1.75	3.0	1.25	Semester
D8	Engineering Foundations (AI amplifiers)	1.09	3.0	1.91	Next quarter
D9	Measurement, Value and AI FinOps	1.61	3.0	1.39	Semester

- Re-assess at the end of each horizon with the same form, then compare rounds with make compare BEFORE=... AFTER=... (the comparison PDF shows the deltas).
- Where available, confirm D4 and D9 with Copilot usage metrics [6] and delivery metrics (D9-Q2) before reporting progress.

7. First 90 days

7.1 Weeks 1 to 4

1. W1: Champions Kickoff
2. W2: Copilot Spaces (shared context)
3. W3: A2A (Agent-to-Agent) protocol
4. W4: Microsoft Foundry overview
5. W4: Office hours #1

7.2 Weeks 5 to 8

- Start S2 Foundry + SRE: Start with D8-Q7 (Automated testing as a control system). Target the L3 anchor: Most services have layered automated tests that run on every PR within agreed time budgets.
- Start S5 GitHub Copilot Acceleration: Start with D9-Q3 (Pull request flow metrics). Target the L3 anchor: PR lifecycle metrics are reported per organization; agent-authored PRs are identified separately, including their post-merge fix rate.

7.3 Weeks 9 to 12

1. W10: Biweekly office hours

8. References

- [1] DORA. DORA AI Capabilities Model. Google Cloud, 2025. <https://dora.dev/ai/capabilities-model/report/>
- [2] DORA. DORA AI Capabilities Model: Survey Questions. Google Cloud, 2025. <https://dora.dev/ai/capabilities-model/questions/>
- [3] Google Cloud. Announcing the 2025 DORA Report: State of AI-Assisted Software Development. 2025. <https://cloud.google.com/blog/products/ai-machine-learning/announcing-the-2025-dora-report>
- [6] GitHub Docs. GitHub Copilot usage metrics. <https://docs.github.com/en/copilot/concepts/billing-and-usage/copilot-usage-metrics/copilot-metrics>
- [7] GitHub Docs. Risks and mitigations for GitHub Copilot cloud agent. <https://docs.github.com/en/enterprise-cloud@latest/copilot/concepts/security-governance-and-network-settings/risks-and-mitigations>
- [8] GitHub Docs. Best practices for using GitHub Copilot to work on tasks. <https://docs.github.com/en/enterprise-cloud@latest/copilot/tutorials/cloud-agent/get-the-best-results>
- [12] GitHub Docs. About GitHub Copilot code review. <https://docs.github.com/en/copilot/concepts/agents/code-review>
- [14] GitHub Docs. Application card: GitHub Copilot Agents. <https://docs.github.com/en/copilot/responsible-use/agents>
- [15] GitHub Docs. Managing policies and features for GitHub Copilot in your organization. <https://docs.github.com/en/copilot/how-tos/administer-copilot/manage-for-organization/manage-policies>
- [16] Microsoft Azure Blog. Agentic DevOps: Evolving software development with GitHub Copilot and Microsoft Azure. 2025. <https://azure.microsoft.com/en-us/blog/agentic-devops-evolving-software-development-with-github-copilot-and-microsoft-azure/>
- [17] Microsoft for Developers. Agentic DevOps in action: Reimagining every phase of the developer lifecycle. 2025. <https://developer.microsoft.com/blog/reimagining-every-phase-of-the-developer-lifecycle/>
- [18] Microsoft Learn. AI strategy: Cloud Adoption Framework. <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/ai/strategy>
- [19] Microsoft Learn. Govern and secure AI agents: Cloud Adoption Framework. <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/ai-agents/governance-security-across-organization>
- [20] Microsoft Learn. Responsible AI policies: Cloud Adoption Framework. <https://learn.microsoft.com/en-us/azure/cloud-adoption-framework/ai/responsible-ai-policies>
- [21] Microsoft. Responsible AI Principles and Approach. <https://www.microsoft.com/en-us/ai/principles-and-approach>
- [22] Microsoft Azure Blog. Announcing Azure Copilot agents and AI infrastructure innovations. 2025. <https://azure.microsoft.com/en-us/blog/announcing-azure-copilot-agents-and-ai-infrastructure-innovations/>
- [23] Anthropic. Anthropic Economic Index: AI's impact on software development. 2025-04-28. <https://www.anthropic.com/research/impact-software-development>
- [24] Anthropic. How Claude Code is used in practice. 2026-06-16. <https://www.anthropic.com/research/claude-code-expertise>
- [25] Anthropic. Effective context engineering for AI agents. 2025-09-29. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>
- [26] Anthropic. Building effective agents. 2024-12-19. <https://www.anthropic.com/engineering/building-effective-agents>
- [27] Anthropic. Best practices for Claude Code. Claude Code Docs. <https://code.claude.com/docs/en/best-practices>
- [28] Gartner. Gartner Says 75% of Enterprise Software Engineers Will Use AI Code Assistants by 2028. 2024-04-11. <https://www.gartner.com/en/newsroom/press-releases/2024-04-11-gartner-says-75-percent-of-enterprise-software-engineers-will-use-ai-code-assistants-by-2028>
- [30] Gartner. Gartner Says Generative AI will Require 80% of Engineering Workforce to Upskill Through 2027. 2024-10-03. <https://www.gartner.com/en/newsroom/press-releases/2024-10-03-gartner-says-generative-ai-will-require-80-percent-of-engineering-workforce-to-upskill-through-2027>

- [31] Gartner. Gartner Says the Market for Enterprise AI Coding Agents Is Entering a New Phase of Expansion and Competitive Realignment. 2026-05-20. <https://www.gartner.com/en/newsroom/press-releases/2026-05-20-gartner-says-the-market-for-enterprise-ai-coding-agents-is-entering-a-new-phase-of-expansion-and-competitive-realignment>
- [32] Gartner. Gartner Predicts AI Coding Costs Will Surpass Average Developer's Salary by 2028 as Token Consumption Surges. 2026-06-24. <https://www.gartner.com/en/newsroom/press-releases/2026-06-24-gartner-predicts-ai-coding-costs-will-surpass-average-developer-salary-by-2028-as-token-consumption-surges>
- [34] Cui, Z. K., Demirer, M., Jaffe, S., Musolff, L., Peng, S., Salz, T. The Effects of Generative AI on High-Skilled Work: Evidence from Three Field Experiments with Software Developers. Management Science, published online 2026-02-27. <https://doi.org/10.1287/mnsc.2025.00535> (pre-print: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4945566)
- [35] METR. Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity. 2025-07-10. <https://metr.org/blog/2025-07-10-early-2025-ai-experienced-os-dev-study/> (paper: <https://arxiv.org/abs/2507.09089>)
- [37] Stack Overflow. 2025 Developer Survey: AI. <https://survey.stackoverflow.co/2025/ai>
- [38] OWASP GenAI Security Project. OWASP Top 10 for LLM Applications 2025. <https://genai.owasp.org/resource/owasp-top-10-for-llm-applications-2025/>
- [39] OWASP GenAI Security Project. OWASP Top 10 for Agentic Applications for 2026. 2025-12-09. <https://genai.owasp.org/resource/owasp-top-10-for-agentic-applications-for-2026/>
- [40] NIST. SP 800-218A: Secure Software Development Practices for Generative AI and Dual-Use Foundation Models: An SSDF Community Profile. 2024-07. <https://csrc.nist.gov/pubs/sp/800/218/a/final>
- [41] NIST. AI Risk Management Framework. <https://www.nist.gov/itl/ai-risk-management-framework>
- [42] ISO. ISO/IEC 42001:2023: AI management systems. <https://www.iso.org/standard/42001>
- [43] OpenSSF. SLSA: Supply-chain Levels for Software Artifacts. <https://slsa.dev/>
- [44] Forsgren, N., Storey, M.-A., Maddila, C., Zimmermann, T., Houck, B., Butler, J. The SPACE of Developer Productivity. ACM Queue 19(1), 2021-03-06. <https://queue.acm.org/detail.cfm?id=3454124> (DOI: <https://doi.org/10.1145/3454122.3454124>)
- [45] Liu, B., Qiu, H., Goiri, Í., Fonseca, R., Bianchini, R., Choukse, E. Agentic Coding in the Wild: Characterizing GitHub Copilot Traces at Production Scale. arXiv:2608.00101, 2026-07-30. <https://arxiv.org/abs/2608.00101>
- [46] Arabat, A., Sayagh, M. Toward Instructions-as-Code: Understanding the Impact of Instruction Files on Agentic Pull Requests. arXiv:2606.13449, 2026-06-11. <https://arxiv.org/abs/2606.13449>
- [47] Denisov-Blanch, Y., Agarwal, S., Azaletskiy, P., He, H., Schaeffer, R., Miranda, B., Vasilescu, B., Koyejo, S. A Few Pages of Markdown: Committed AI Configuration and Lower Quality Cost after Coding-Agent Adoption. arXiv:2608.25241, 2026-08-26. <https://arxiv.org/abs/2608.25241>
- [48] Shen, J. H., Tamkin, A. How AI Impacts Skill Formation. arXiv:2601.20245, 2026-01-28. <https://arxiv.org/abs/2601.20245> (Anthropic summary: <https://www.anthropic.com/research/AI-assistance-coding-skills>)
- [49] Pinna, G., Gong, J., Williams, D., Sarro, F. Comparing AI Coding Agents: A Task-Stratified Analysis of Pull Request Acceptance. arXiv:2602.08915, 2026-02-09. <https://arxiv.org/abs/2602.08915>
- [50] Takerngsaksiri, W., Duong, N., Barnett, S. Who Finishes the Job? A Study of Follow-Up Fixes and Commit Authorship on AI Coding Agent Pull Requests. arXiv:2609.26847, 2026-09-22. <https://arxiv.org/abs/2609.26847>
- [52] Sakib, A. H. M. N., Banik, D., Jadliwala, M. Trust but Verify? Uncovering the Security Debt of Autonomous Coding Agents. arXiv:2607.12428, 2026-07-14. <https://arxiv.org/abs/2607.12428>
- [54] Selvanayagam, N., Ghaleb, T. A. AI-to-AI Code Reviews of GitHub Pull Requests. arXiv:2608.21311, 2026-08-21. <https://arxiv.org/abs/2608.21311>
- [55] Stolze, M., Strässle, M. When Review Alone No Longer Scales: Layered Supervision in AI-Assisted Software Engineering. arXiv:2608.26316, 2026-08-26. <https://arxiv.org/abs/2608.26316>
- [56] Monperrus, M. The End of Code Review: Coding Agents Supersede Human Inspection. arXiv:2606.13175, 2026-06-11. <https://arxiv.org/abs/2606.13175>
- [58] Farrag, S. E. The Productivity-Reliability Paradox: Specification-Driven Governance for AI-Augmented Software Development. arXiv:2605.01160, 2026-05-01. <https://arxiv.org/abs/2605.01160>